Zeek + Filebeat + Elasticsearch (ELK) Deployment Report

---

1. Introduction (Purpose of the Project)

The objective of this project is to deploy a complete network monitoring and SIEM pipeline using **Zeek, Filebeat, Elasticsearch, and Kibana**.

This setup is designed specifically for **Security Operations Center (SOC)** use cases such as:

• Network traffic monitoring
• Suspicious activity detection
• Log analysis and investigation
• Centralized visibility of network events

This project simulates a **real-world SOC environment**, where logs are continuously collected, processed, stored, and visualized for effective security analysis and incident investigation.

---

2. Why We Use Only These Software (SOC Perspective)

2.1 Why Zeek

Zeek is a **Network Security Monitor**, not just a packet sniffer.

Key reasons for using Zeek:
• Converts raw network traffic into structured security logs
• Produces human-readable logs such as:

- conn.log – connection metadata
- dns.log – DNS activity
- http.log – HTTP request details

SOC analysts work with **events**, not raw packets, making Zeek ideal for SOC environments.

2.2 Why Not Wireshark

• Wireshark is packet-based and requires manual analysis
• Not suitable for continuous, large-scale SOC monitoring
• No centralized log management

2.3 Why Filebeat

Filebeat is used for **secure and reliable log shipping**.

Advantages:
• Automatically reads Zeek log files
• Handles log rotation efficiently
• Prevents data loss
• Provides a built-in Zeek module
• Industry-standard log shipper for SOC environments

2.4 Why Not Custom Scripts

• Custom scripts may crash or fail silently
• No retry or fault-tolerance mechanism
• Not considered SOC-standard or enterprise-ready

2.5 Why Elasticsearch

Elasticsearch acts as the **SIEM data storage and search engine**.

Benefits:
• Stores massive volumes of security logs
• High-speed indexing and searching
• Optimized for time-based log analysis

2.6 Why Not SQL Databases

• SQL databases are not optimized for log analytics
• Slower querying for time-series security data
• Poor scalability for SOC workloads

2.7 Why Kibana

Kibana serves as the **SOC analyst interface**.

Features:
• Interactive dashboards
• Advanced searching and filtering
• Timeline-based investigations
• No command-line requirement for analysts

---

3. Architecture Overview (SOC Data Flow)

Network Traffic
↓
Zeek
↓
Filebeat
↓
Elasticsearch
↓
Kibana

This architecture represents a **standard SOC / SIEM pipeline** used in enterprise environments.

---

4. System Environment

Operating System: Ubuntu 24.04 LTS
Zeek Version: 8.0.4
Filebeat Version: 7.17.29 / 8.19.8
Elasticsearch Version: 7.17.29

Kibana Version: 7.17.29
Network Interface: ens33

---

5. Step-by-Step Deployment with Explanation

Step 1: Identify Network Interface

Command:

ip a

Reason:
Zeek must listen on the correct network interface.
Incorrect interface selection prevents packet capture.

Step 2: Configure Zeek Interface

Configuration File:

/opt/zeek/etc/node.cfg

Change:

interface=eth0

To:

interface=ens33

Reason:
Ensures Zeek captures live network traffic and avoids startup errors.

Step 3: Deploy Zeek

Commands:

sudo /opt/zeek/bin/zeekctl deploy

sudo /opt/zeek/bin/zeekctl status

Reason:
• Starts Zeek services
• Applies configuration changes
• Confirms Zeek is running successfully

Step 4: Verify Zeek Logs

Command:

sudo ls /opt/zeek/logs/current

Reason:
SOC monitoring depends entirely on log generation.
No logs means no security visibility.

---

6. Issues Faced and Resolution

Issue 1: Zeek Interface Error

Error:

pcap_error: No such device exists

Cause:
Incorrect network interface configured.

Solution:
• Updated interface name in node.cfg
• Redeployed Zeek

Issue 2: Zeek Logs Not Accessible

Error:

Permission denied

Cause:
• Zeek logs owned by root
• Filebeat runs as a non-root user

Solution Commands:

sudo useradd filebeat

sudo chgrp -R filebeat /opt/zeek/logs

sudo chmod -R 750 /opt/zeek/logs

Reason:
• Follows principle of least privilege
• Allows Filebeat to safely read logs

Issue 3: JSON Parsing Error

Error:

json: cannot unmarshal number into Go value

Cause:
Zeek logs are generated in **TSV format**, not JSON.

Solution:
• Enabled Zeek module in Filebeat
• Disabled manual JSON parsing

Configuration File:

/etc/filebeat/modules.d/zeek.yml

Issue 4: Kibana Dashboard Import Error

Error:

Saved objects are not backwards compatible

Cause:
Filebeat version newer than Kibana.

Resolution:
• Skipped dashboard import
• Used Kibana Discover for analysis

Reason:
Log ingestion was successful; visualization through Discover was sufficient.

---

7. Filebeat Configuration and Start

Commands:

sudo filebeat test config

sudo systemctl restart filebeat

sudo systemctl status filebeat

Reason:
• Validates configuration
• Starts log shipping to Elasticsearch
• Confirms Filebeat service health

---

8. Elasticsearch Verification

Service Check: curl http://localhost:9200

Reason:
Confirms Elasticsearch is running.

Index Verification:

curl http://localhost:9200/_cat/indices/filebeat*?v

Reason:
Confirms logs are indexed successfully.
This proves SIEM ingestion is working.

---

9. Kibana Access and SOC Validation

Access URL: http://localhost:5601

Discover Configuration:
• Index pattern: filebeat-*
• Filter:

event.module : zeek

Reason:
SOC analysts validate logs through Kibana to ensure data is searchable and usable.

---

10. SOC Analysis Use Case

With this deployment, a SOC analyst can:

• Detect suspicious DNS queries
• Analyze abnormal network connections
• Investigate HTTP traffic
• Perform timeline-based incident analysis

---

11. Final Deployment Checklist

Zeek capturing traffic – ✓
Logs generated – ✓
Filebeat shipping logs – ✓
Elasticsearch indexing data – ✓
Kibana visualizing logs – ✓

---

12. Conclusion

The **Zeek–Filebeat–ELK Stack** was successfully deployed for SOC analysis.
All issues related to network interface configuration, permissions, log format handling, and version compatibility were resolved through systematic troubleshooting.

This deployment represents a **real-world SOC / SIEM monitoring architecture**, making it highly suitable for **security internships, academic projects, and practical SOC operations**.
